

CUDA STREAMS AND EVENTS

25. APRIL 2018 | JOCHEN KREUTZ

OVERVIEW

- Manual memory management
- Pinned (pagelocked) host memory
- Asynchronous and concurrent memory copies
- CUDA streams
 - The default stream and the `cudaStreamNonBlocking` flag
- CUDA Events
- CUBLAS
- nvprof + nvvp recap

GETTING DATA IN AND OUT

- GPU has separate memory
- Allocate memory on device
- Transfer data from host to device
- Transfer data from device to host
- Free device memory

GETTING DATA IN AND OUT

Allocate device memory

```
cudaMalloc ( T** pointer, size_t nbytes )
```

Example:

```
// Allocate a vector of 2048 floats on device
```

```
float * a_gpu;
```

```
int n = 2048;
```

```
cudaMalloc ( &a_gpu, n * sizeof(float) );
```

↑
Address of pointer

↖
Get size of a float

COPY FROM HOST TO DEVICE

Copy data from host to device memory

```
cudaMemcpy(void* dst, void* src, size_t nbytes,  
            enum cudaMemcpyKind dir)
```

Example:

```
// Copy vector of floats a of length n=2048  
// to a_gpu on device  
cudaMemcpy(a_gpu, a, n * sizeof(float),  
           cudaMemcpyHostToDevice);
```

COPY FROM DEVICE TO HOST

Copy data from host to device memory

```
cudaMemcpy(void* dst, void* src, size_t nbytes,  
            enum cudaMemcpyKind dir)
```

Example:

```
// Copy vector of floats a_gpu of length  
// n=2048 to a on host  
cudaMemcpy(a, a_gpu, n * sizeof(float),  
           cudaMemcpyDeviceToHost);
```

GETTING DATA IN AND OUT

- Allocate memory on device

```
cudaMalloc(void** pointer, size_t nbytes)
```

- Transfer data between host and device

```
cudaMemcpy(void* dst, void* src, size_t nbytes,  
            enum cudaMemcpyKind dir)  
//dir is cudaMemcpyHostToDevice or cudaMemcpyDeviceToHost
```



- Free device memory

```
cudaFree(void* pointer)
```

PINNED HOST MEMORY

- Host memory allocated with **malloc** is pagable
 - Memory pages associated with the memory can be moved around by the OS Kernel, e.g. to swap space on hard disk
- Transfers to and from the GPU memory need to go over PCI-E
 - PCI-E transfers are handled by DMA engines on the GPU and work independently of the CPU/OS kernel
 - If OS kernel moves memory pages involved in such a DMA transfer the wrong data will be moved
 - Pinning memory pages inhibit the OS kernel from moving them around and make them usable to DMA transfer


PINNED HOST MEMORY



Host memory allocated with **malloc** is staged through a pinned memory buffer managed by the CUDA Driver

- No asynchronous memory copies are possible (CPU interaction is necessary to drive the pipeline)
- Higher latency and lower bandwidth compared to DMA transfers

PINNED HOST MEMORY



memcpy D->H



Time

Using pinned host memory

- enables asynchronous memory copies
- Lowers latency and increases bandwidth

PINNED HOST MEMORY – HOW TO USE IT?

- Using POSIX functions like **mlock** is not sufficient, because the CUDA driver needs to know that the memory is pinned
- Two ways to get pinned host memory
 - Using **cudaMallocHost/cudaFreeHost** to allocate new memory
 - Using **cudaHostRegister/cudaHostUnregister** to pin memory after allocation
- **cudaMemcpy** makes automatic use of it
- **cudaMemcpyAsync** can be used to issue asynchronous memory copies
- Can be directly accessed from Kernels (zero-copy) – use **cudaHostGetDevicePointer**

CUDA STREAMS

- CUDA Streams are work queues to express concurrency between different tasks, e.g.
 - host to device memory copies
 - device to host memory copies
 - kernel execution
- To overlap different tasks just launch them in different streams
 - All tasks launched into the same stream are executed in order
 - Tasks launched into different streams might execute concurrently (depending on available resources: two copy engines, compute resources)

CUDA STREAMS – HOW TO USE THEM?

- Create/Destroy

```
cudaStream_t stream;  
cudaStreamCreate ( &stream );  
cudaStreamDestroy ( stream );
```

- Launch

```
my_kernel<<<grid,block,0,stream>>>(...);  
cudaMemcpyAsync( ..., stream );
```

- Synchronize


```
cudaStreamSynchronize( stream );
```

CUDA STREAMS – THE DEFAULT (NULL) STREAM

- Kernel launches are always asynchronous
 - Which stream is used here ?


```
my_kernel<<<grid,block>>> (...);
```

- The default (null) stream



CUDA STREAMS – THE DEFAULT (NULL) STREAM

- **Watch out for false dependencies !**
 - The default stream waits for work in all other streams which do not have the **cudaStreamNonBlock** flag set



User streams with the **cudaStreamNonBlocking** flag set can execute concurrently to stream 0

CUDA EVENTS

CUDA Events are synchronization markers that can be used to:

- Time asynchronous tasks in streams
- Allow fine grained synchronization within a stream
- Allow inter stream synchronization, e.g. let a stream wait for an event in another stream

CUDA EVENTS – HOW TO USE THEM?

Create/Destroy

```
cudaEventCreate( &event );  
cudaEventDestroy( event );
```

Record

```
cudaEventRecord( event, stream );
```

Query

```
cudaEventQuery( event );
```

Synchronize

```
cudaEventSynchronize( event );
```

Timing

```
cudaEventElapsedTime( &time, start, end );
```

CUDA EVENTS – EXAMPLE FOR KERNEL TIMING

KERNEL TIMING

```
cudaEventRecord ( startEvent, stream );  
my_kernel<<<grid,block,0,stream>>>(...);  
cudaEventRecord ( endEvent, stream );  
  
//Host can do other work  
  
//Get runtime of my_kernel in ms  
float runtime = 0.0f;  
cudaEventSynchronize ( endEvent );  
cudaEventElapsedTime ( &runtime, startEvent, endEvent );
```

CALLING CUBLAS

HOW TO USE CUBLAS FUNCTION

```
#include "cublas_v2.h"
...
cublasHandle_t handle;

//Initialize cuBLAS
cublasCreate(&handle);

//Set cuBLAS execution stream
cublasSetStream(handle, stream);

//Call SAXPY
cublasSaxpy(handle, n, &alpha, x, 1, y, 1);
...
//Free up resources
cublasDestroy(handle);
```

THE COMMAND LINE PROFILER NVPROF

- Simple launcher to get profiles of your application
- Profiles CUDA Kernels and API calls

HOW TO USE COMMAND LINE PROFILER

```
> nvprof ./jacobi
```

```
===== NVPROF is profiling jacobi...
```

```
===== Command: jacobi
```

```
Jacobi (serial)
```

```
[...] snip
```


```
===== Profiling result:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
72.14	352.65ms	1000	352.65us	350.48us	354.94us	Jacobi_86_gpu
26.02	127.23ms	1000	127.23us	93.48us	128.34us	Jacobi_74_gpu
0.84	4.09ms	1000	4.09us	4.04us	4.36us	Jacobi_96_gpu_red
0.61	3.00ms	1009	2.97us	2.78us	56.16us	[CUDA memcpy HtoD]
0.39	1.91ms	1002	1.91us	1.82us	52.41us	[CUDA memcpy DtoH]

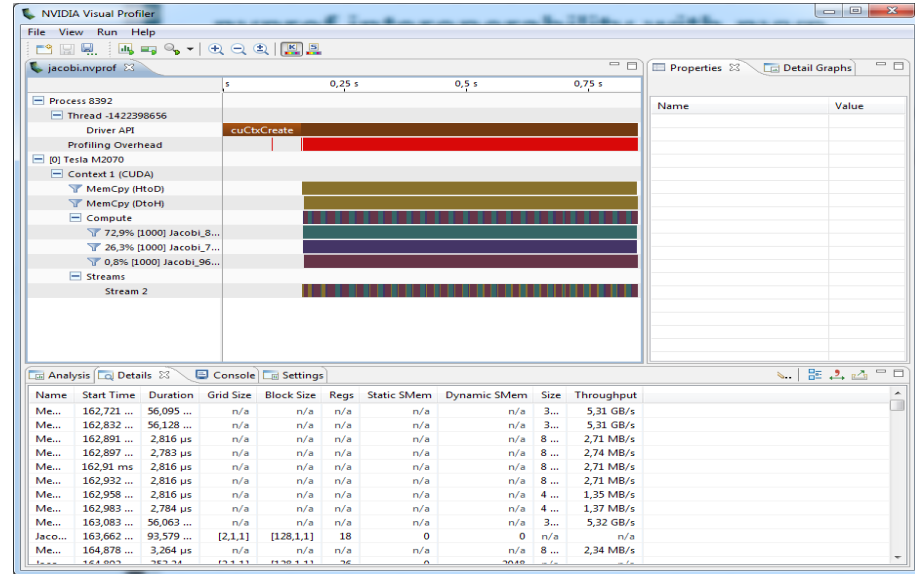
NVPROF INTEROPERABILITY WITH NVVP

- nvprof can write the application profile to nvvp compatible file:

```
nvprof -o jacobi.nvprof ./jacobi
```



Import with nvvp



NVPROF IMPORTANT COMMAND-LINE OPTIONS

HOW TO USE COMMAND LINE PROFILER

Options:

`-o, --output-profile <filename>`

Output the result file which can be imported later or opened by the NVIDIA Visual Profiler.

`--events <event names>`

Specify the events to be profiled on certain device(s). Multiple event names separated by comma can be specified. Which device(s) are profiled is controlled by the '--devices' option. Otherwise events will be collected on all devices.

For a list of available events, use '--query-events'.

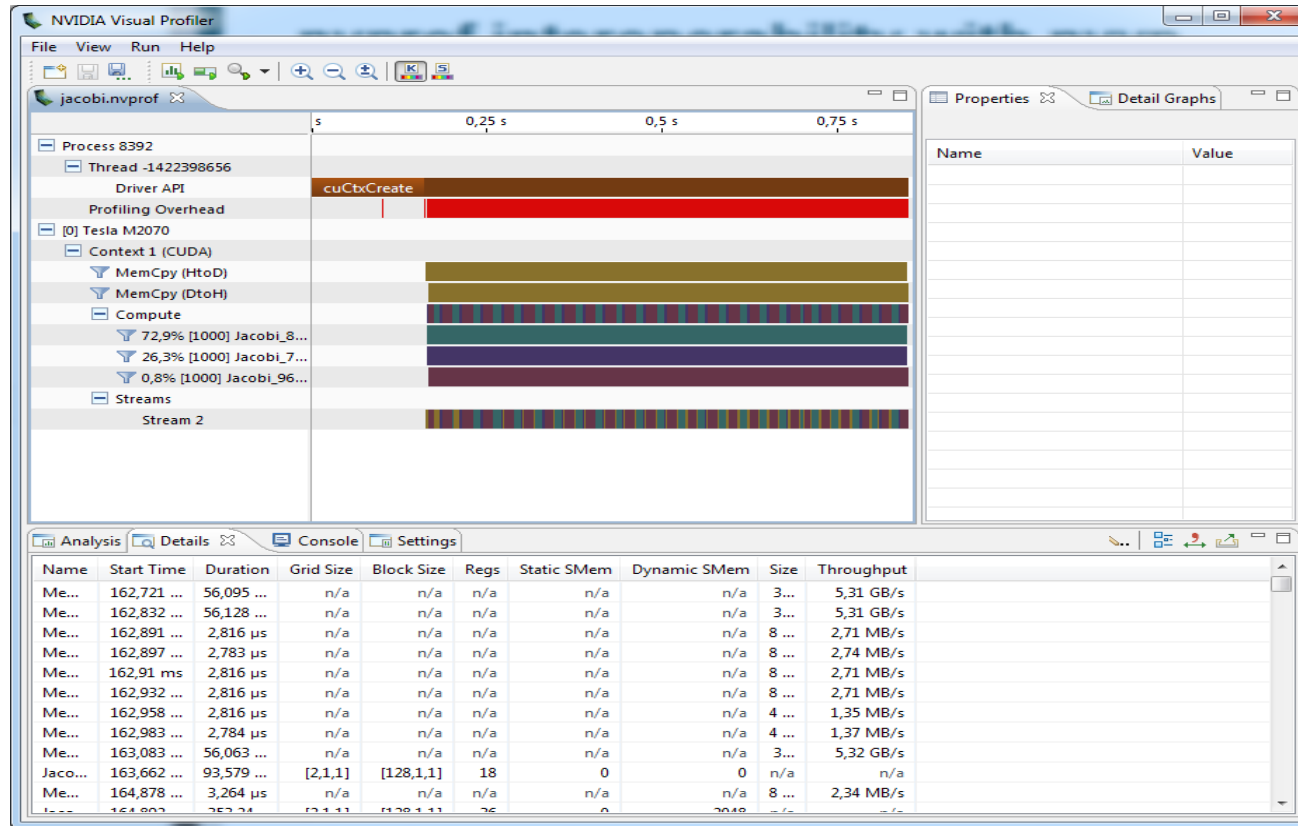
`--query-events`

List all the events available on each device.

`-h, --help`

Print this help information.

VISUAL PROFILE - NVVP



EXERCISES: TASK 1

Optimize and overlap host to device and device to host transfers

- Task 1a:
 - Follow TODOs in ***CUDAStreams/exercises/tasks/task1a.cu***
 - Allocate host buffers in pinned memory
 - View nvprof profile in nvvp
- Task 1b:
 - Follow TODOs in ***CUDAStreams/exercises/tasks/task1b.cu***
 - Create Upload and Download Stream
 - Issue Host to Device and Device to Host Transfer asynchronously in the two new streams.
 - View nvprof profile in nvvp
-

EXERCISES: TASK 2

- Follow TODOs in ***CUDAStreams/exercises/tasks/task2.cu***
 - Set CUBLAS execution stream
 - Call CUBLAS SAXPY
 - Fix position of cudaStreamSynchronize
- Instructions can be found in ***Instructions.rst***
- Solutions in ***CUDAStreams/exercises/solutions***
- Slides can be found in in
CUDAStreams/slides/CUDAStreams_and_Events.pdf